

Chat with PDF Project

Absolutely! Here's a **detailed, hierarchical syllabus of Generative AI** — broken down into **core modules, sub-topics, and leaf-level topics**. This is designed to take you from foundational AI concepts to building advanced GenAI apps like ChatGPT-like tools, image generators, and retrieval-augmented systems.

COMPLETE GENERATIVE AI SYLLABUS — DETAILED LEAF TOPICS

Module 1: Foundations of AI, ML, and Deep Learning

1.1 Introduction to AI

- What is Artificial Intelligence (AI)?
- Types of AI: Narrow, General, Super
- AI vs ML vs DL vs GenAI
- Use Cases of AI

1.2 Machine Learning Basics

- Supervised, Unsupervised, Reinforcement Learning
- Regression and Classification
- Decision Trees, SVMs, KNN

1.3 Deep Learning Essentials

- Neural Networks (ANN)
- Forward and Backward Propagation
- Activation Functions (ReLU, Sigmoid, Tanh, Softmax)
- Loss Functions (Cross-Entropy, MSE)

- Optimization Algorithms (SGD, Adam)

1.4 Tools & Frameworks

- Python, NumPy, Pandas, Matplotlib
 - scikit-learn, TensorFlow, PyTorch basics
-



Module 2: Transformers & Foundation Models

2.1 Understanding Transformers

- Sequence Models (RNN, LSTM overview)
- Attention Mechanism
- Scaled Dot-Product Attention
- Self-Attention vs Cross-Attention
- Positional Encoding
- Encoder-Decoder Architecture

2.2 Modern Transformers

- BERT, RoBERTa
 - GPT-1, 2, 3, 3.5, 4
 - T5, BART, LLaMA, Mistral, Claude
 - Vision Transformers (ViT)
 - Instruction tuning vs Fine-tuning vs RLHF
-



Module 3: NLP with Large Language Models (LLMs)

3.1 Natural Language Processing

- Tokenization (Byte-Pair Encoding, WordPiece)
- Stopwords, Lemmatization, Stemming

- NER, POS tagging, Chunking
- Sentence and Document Embeddings

3.2 Large Language Models

- Architecture of GPT & LLaMA
- Pre-training and Transfer Learning
- In-context Learning & Few-shot Learning
- Prompt Engineering (detailed in Module 4)

3.3 Hands-On with APIs

- Using OpenAI GPT API (chat/completions)
 - Using HuggingFace Transformers
 - Model inference with `pipeline()` & `AutoModel`
-

Module 4: Prompt Engineering (PE)

4.1 Prompt Design Techniques

- Zero-shot Prompting
- One-shot and Few-shot Prompting
- Chain of Thought (CoT)
- Role-based Prompting

4.2 Prompt Engineering Tools

- OpenAI Playground
- LangChain PromptTemplates
- LlamaIndex Prompt Frameworks

4.3 Prompt Injection & Safety

- Jailbreak Prompts

- Preventing Prompt Injection
 - Guardrails & Moderation APIs
-

Module 5: Embeddings & Vector Databases

5.1 Embedding Concepts

- What are Embeddings?
- Word2Vec, GloVe, FastText
- SentenceTransformers & OpenAI Embeddings

5.2 Vector Similarity

- Cosine Similarity
- Euclidean Distance
- Approximate Nearest Neighbor Search (ANN)

5.3 Vector Databases

- FAISS: IndexFlatL2, IndexIVFFlat
 - ChromaDB
 - Pinecone
 - Weaviate, Qdrant (intro only)
-

Module 6: LangChain & LlamaIndex

6.1 LangChain Core Concepts

- Documents and Loaders
- Text Splitters
- Embedding & Vector Store
- RetrievalQA Chain

- Conversational Memory
- Agents and Tool Usage

6.2 LlamaIndex Essentials

- Document ingestion & chunking
 - Indexing documents
 - Query Engines (Simple, Retriever, Composable Graphs)
 - Response Synthesizers
-



Module 7: Retrieval-Augmented Generation (RAG)

7.1 RAG Pipeline

- Extracting + Chunking PDFs, DOCs
- Generating and storing embeddings
- Query + Top-k Retrieval
- Building the Prompt + LLM Response

7.2 Applications of RAG

- Chat with PDF
 - Chat with YouTube
 - Internal Docs Chatbot
 - Legal or Medical AI Assistant
-



Module 8: Multimodal GenAI (Text, Image, Audio, Video)

8.1 Text-to-Image Generation

- DALL·E, Midjourney, Stable Diffusion
- Prompting for Image Generation

- ControlNet, Inpainting

8.2 Image-to-Text & Vision Models

- BLIP, GPT-4o, Gemini Vision
- Visual Question Answering (VQA)
- CLIP (Contrastive Language-Image Pretraining)

8.3 Audio + Video

- Whisper for Speech-to-Text
 - Bark / Tortoise / ElevenLabs for Text-to-Speech
 - RunwayML for Video Gen
 - Sora (OpenAI) — Concepts only
-



Module 9: Model Fine-tuning & Training

9.1 Fine-Tuning Fundamentals

- Full fine-tuning vs LoRA
- PEFT: Parameter Efficient Fine Tuning
- Dataset preparation for training

9.2 Tools for Training

- Hugging Face `transformers`, `datasets`, `accelerate`
- `peft`, `bitsandbytes`, `transformers` Trainer

9.3 Evaluating Models

- BLEU, ROUGE, METEOR (for text gen)
 - Embedding similarity evaluation
 - Human evaluation techniques
-



Module 10: LLM App Development & Deployment

10.1 Backend Development

- FastAPI, Flask basics
- RESTful API creation for GenAI
- Upload endpoints, file parsing

10.2 Frontend Development

- React.js, Tailwind CSS
- shadcn/ui, file upload components
- Gradio / Streamlit for Python-based UI

10.3 Full-stack Integration

- Connecting React frontend to FastAPI
- Hosting Vector DBs (locally or on cloud)
- Fetch + display LLM responses

10.4 Deployment

- Vercel (frontend)
 - Render / Railway (backend)
 - Hugging Face Spaces (streamlit/gradio)
 - Docker (optional)
-



Module 11: Ethics, Safety, and Responsible GenAI

11.1 Ethical Considerations

- Bias in Language Models
- Toxicity, Misinformation

11.2 Security









- Prompt Injection Attacks
- Data Privacy in GenAI

11.3 Governance & Fair Use

- Model licensing (Open vs Closed)
 - Copyright & Fair Use in Generated Content
-



Module 12: Capstone Projects (Build & Deploy)

-  Chat with PDF (RAG + LangChain + Chroma)
 -  Chat with YouTube Transcript
 -  AI Voice Assistant (Speech-to-Text + LLM + TTS)
 -  AI Resume Analyzer / Generator
 -  AI Image Generator from Prompt
 -  Semantic Search Engine
 -  Medical Q&A Chatbot
 -  AI Job Interviewer (Q&A)
-

Would you like this as a downloadable **Notion template**, **PDF**, or **Excel planner** to track your GenAI journey?